

# Caching & Content Delivery Network (CDN)

## I. The Role of Caching

Caching is the first line of defense in System Design. It acknowledges a simple truth: Fetching data from memory (RAM) is orders of magnitude faster than fetching it from a disk (Database).

### 1. Where do we Cache?

- Client-side: Storing assets in the browser to avoid redundant network calls.
- CDN (The Edge): Storing static assets (images, JS, CSS) at servers physically close to the user.
- Application Cache: Using tools like Redis or Memcached to store database query results or session data.

Interview Talk Track:

"I implement caching to solve two main problems: High Latency and Database Bottlenecks. By moving frequently accessed, 'read-heavy' data from the disk into RAM, we can reduce response times from 100ms+ down to sub-10ms. This not only makes the app feel faster but also saves our database from crashing during traffic spikes."

---

## II. Caching Strategies (Write/Read Patterns)

How you move data into the cache is a critical architectural decision.

### 1. Cache-Aside (Lazy Loading)

The application looks at the cache first. If it's a miss, it fetches from the DB and updates the cache.

- Best for: General-purpose, read-heavy workloads.

### 2. Write-Through

Data is written to the cache and the DB at the same time.

- Best for: Maintaining high consistency between cache and DB.

### 3. Write-Behind (Write-Back)

The app writes only to the cache. The cache then updates the DB asynchronously (later).

- Best for: Write-heavy applications (like logging or real-time counters).
- Risk: If the cache server crashes before the DB is updated, data is lost.

Interview Talk Track:

"When deciding on a caching strategy, I look at the Read vs. Write ratio. For most web apps, I use the Cache-Aside pattern because it's simple and robust. However, if the system needs to handle massive bursts of writes—like a 'Like' button on a viral post—I'd consider a Write-Behind strategy to batch those updates and reduce the pressure on my database."

---

### III. Cache Invalidation & Eviction

"There are only two hard things in Computer Science: cache invalidation and naming things." — *Phil Karlton*

#### 1. Invalidation: How to handle stale data?

- TTL (Time to Live): Set an expiration timer.
- Write-through: Update cache whenever the DB is updated.

#### 2. Eviction: What happens when the cache is full?

- LRU (Least Recently Used): Discards the data that hasn't been looked at for the longest time. (Most common).
  - LFU (Least Frequently Used): Discards data that is accessed the least often.
- 

### IV. Content Delivery Network (CDN)

A CDN is essentially a distributed cache for the entire planet.

- Edge Servers: Located in "Points of Presence" (PoPs) globally.
- Origin Server: Your main backend server.

Interview Talk Track:

"A CDN is essential for any global application. Without it, a user in Tokyo would have to wait for a 200ms round-trip to a server in New York just to load a logo. By using a CDN, we serve that logo from a server in Tokyo itself. This

significantly reduces the 'Time to First Byte' (TTFB) and protects our origin server from being hammered by requests for static files."

---

## Practice Interview Questions (Day 17–18)

### 1. "What is the 'Cache Stampede' problem, and how do you fix it?"

- Answer: A Cache Stampede happens when a very popular item expires (TTL ends) and thousands of users hit the database at the same time to refresh it. We fix this by using 'X-Fetch' (re-calculating the value slightly before it expires) or using locking so only one request goes to the DB to refresh the cache.

### 2. "Why shouldn't you cache everything?"

- Answer: Caching adds complexity and costs. If data changes constantly (like a stock price ticker), the overhead of invalidating the cache might make it slower than just reading from the DB. Caching is for data that is read frequently but changed infrequently.

### 3. "Explain the difference between a Pull CDN and a Push CDN."

- Answer: \* Pull CDN: The CDN asks your server for content only when a user requests it. (Lower maintenance).
  - Push CDN: You manually upload content to the CDN. (Better for very large files or infrequent updates).

### 4. "How do you handle sensitive user data (like a bank balance) in a cache?"

- Answer: We generally avoid caching highly sensitive, rapidly changing data. If we must, we ensure the cache is encrypted, has a very short TTL, and is never cached at the 'Edge' (CDN) where it could be exposed to other users.

### 5. "What metric helps you decide if your cache is effective?"

- Answer: The Cache Hit Ratio. If your ratio is 95%, your cache is doing great. If it's 10%, you are likely caching the wrong data or your TTL is too short, and you're still putting too much load on your database.