

# Logging, Monitoring & Observability (Day 21–22)

## I. The "Three Pillars" of Observability

Observability is the ability to understand the internal state of a system based on the external data it produces. It is built on three pillars: Logs, Metrics, and Traces.

Pillar	What it is	Goal
Logs	Discrete events (text or JSON).	"What happened at a specific time?"
Metrics	Aggregated numbers (CPU, error rate).	"Is the system healthy over time?"
Traces	End-to-end journey of a request.	"Where did the delay happen in the flow?"

Interview Talk Track:

"I differentiate between Monitoring and Observability by saying: Monitoring tells me *when* something is wrong (the 'known-knowns'), but Observability allows me to ask *why* something is wrong, even if I've never seen that specific failure before (the 'unknown-unknowns'). In a microservices environment, you need all three pillars to effectively debug cross-service issues."

---

## II. Logging: From Text to Searchable Data

Logging is the historical record of your application's life.

### 1. Structured Logging (JSON)

Instead of logging: User 123 logged in from 1.1.1.1,

We log: `{"event": "login", "user_id": 123, "ip": "1.1.1.1", "status": "success"}`.

- Why? You can query it like a database. You can ask: *"Give me all failed logins for User 123 in the last 10 minutes."*

## 2. Centralized Logging (The ELK Stack)

In a system with 100 servers, you cannot SSH into each one to read logs. We use the ELK Stack (Elasticsearch, Logstash, Kibana) or Loki to push all logs into one searchable dashboard.

---

## III. Monitoring Frameworks: RED vs. USE

Not all metrics are equal. We use different frameworks depending on what we are looking at.

[Image comparing RED and USE metrics frameworks for system monitoring]

### 1. RED Metrics (For Services/APIs)

- Rate: Number of requests per second.
- Errors: Number of failed requests.
- Duration: Time taken for requests (Latency).

### 2. USE Metrics (For Infrastructure/Hardware)

- Utilization: How busy the resource is (e.g., 80% CPU).
- Saturation: How much extra work is waiting in a queue (e.g., CPU Load Average).
- Errors: Hardware or system-level error counts.

Interview Talk Track:

"When I build dashboards, I follow the RED method for my application services to ensure the user experience is healthy. Simultaneously, I monitor USE metrics for my underlying infrastructure (like DBs and EC2 instances) to detect bottlenecks or resource exhaustion before they impact the application layer."

---

## IV. Distributed Tracing: Connecting the Dots

In a microservices architecture, one user request might hit 5 different services. If the request is slow, which service is the culprit?

- Trace ID: A unique ID generated at the start of the request and passed to every service in the header.
  - Span ID: A unique ID for each individual step (e.g., "Database Query" or "Auth Check") within that trace.
- 

## V. Reliability Targets: SLA, SLO, SLI

- SLI (Indicator): What you measure (e.g., "Successful request rate").
  - SLO (Objective): Your internal target (e.g., "99.9% of requests must succeed").
  - SLA (Agreement): The legal contract (e.g., "If we drop below 99.5%, we pay you back money").
- 

## Practice Interview Questions (Day 21–22)

### 1. "How do you avoid 'Alert Fatigue' in a large-scale system?"

- Answer: Alerts should be actionable. If an alert triggers but an engineer can't do anything about it, it shouldn't be an alert—it should be a dashboard metric. We also set thresholds based on Symptoms (e.g., high error rate) rather than Causes (e.g., high CPU), as CPU might spike naturally without affecting users.

### 2. "Why is JSON logging preferred over plain text in production?"

- Answer: Scalability and Machine Readability. At scale, we use log aggregators like Elasticsearch. JSON allows these tools to index fields automatically, enabling us to perform complex filtering and statistical analysis on logs without writing complex regular expressions (regex).

### 3. "Explain the difference between Latency and Throughput."

- Answer: Latency is the time it takes for a single request to complete (measured in ms). Throughput is the total number of requests the system can handle in a given time (measured in RPS - Requests Per Second). A system can have high throughput but high latency if it's processing many requests slowly.

### 4. "How do you propagate a Trace ID through multiple services?"

- Answer: Through Context Propagation. We use HTTP headers (like X-Trace-ID or W3C Trace Context) to pass the ID. Every time one service calls another, it includes the current Trace ID in the outgoing request headers.

### 5. "What is an 'Error Budget'?"

- Answer: Derived from our SLO. If our SLO is 99.9%, our Error Budget is 0.1%. It represents the amount of downtime or errors we can "afford" in a month. If we haven't used our budget, we can release new features. If we've exhausted it, we stop new releases and focus purely on stability.